

Collective operations

Victor Eijkhout
Jack Gaither

2026 PCSE

Overview

In this section you will learn ‘collective’ operations, that combine information from all processes.

Commands learned:

- *MPI_Bcast, MPI_Reduce, MPI_Gather, MPI_Scatter*
- *MPI_All... variants, MPI_...v variants*
- *MPI_Barrier, MPI_Alltoall, MPI_Scan*

Technically

Routines can be 'collective on a communicator':

- They involve a communicator;
- if one process calls that routine, every process in that communicator needs to call it
- Mostly about combining data, but also opening shared files, declaring 'windows' for one-sided communication.

Collectives





Gathering and spreading information:

- One process has data, you want to spread it around.
- Every process has data, you want to bring it together;

Root process: the one doing the collecting or disseminating.

Basic cases:

- Send the same data to everyone: broadcast.
- Send individual data to each process: scatter.
- Collect data: gather.
- Collect data and compute some overall value (sum, max): reduction.



Exercise 1

How would you realize the following scenarios with MPI collectives?

1. Let each process compute a random number. You want to print the maximum of these numbers to your screen.
2. Each process computes a random number again. Now you want to scale these numbers by their maximum.
3. Let each process compute a random number. You want to print on what processor the maximum value is computed.

Think about time and space complexity of your suggestions.

Allreduce: reduce-to-all


Case 2 in the exercise above contains a common case:
do a reduction, but everyone needs the result.

- *MPI_Allreduce* does the same as:
MPI_Reduce (reduction) followed by *MPI_Bcast* (broadcast)
- Same running time as either, half of reduce-followed-by-broadcast
(no proof given here)
- Common use case, symmetrical expression.

Motivation for allreduce

Example: normalizing a vector

$$y \leftarrow x / \|x\|$$




Norm needs to be known everywhere.
How do you compute it?

Structure of allreduce

- Vectors x, y are distributed: every process has certain elements
- The norm calculation is an all-reduce: every process gets same value
- Every process scales its part of the vector.
- Question: what kind of reduction do you use for an inf-norm?
One-norm? Two-norm?

Conceptual picture

Recall SPMD: every process has the input and output variable



(What actually happens is a different story!)

Exercise 2

How many objections can you come up to this strategy:

1. Gather vector x on some root process;
2. Compute the reduction on that root;
3. Construct the scaled vector on the root;
4. Scatter the scaled vector.

Exercise 3

Standard deviation:

$$\sigma = \sqrt{\frac{1}{N} \sum_i^N (x_i - \mu)^2} \quad \text{where} \quad \mu = \frac{\sum_i^N x_i}{N}$$

and assume that every process stores just one x_i value;
the standard deviation is printed out, and not further used in
computations.

Argue that there is one all-reduce and one reduce with a root.

MPI_Allreduce

Name	Param name	Explanation	C type	F type
MPI_Allreduce (				
MPI_Allreduce_c (				
	sendbuf	starting address of send buffer	const void*	TYPE(*), DIMENSION(..)
	recvbuf	starting address of receive buffer	void*	TYPE(*), DIMENSION(..)
	count	number of elements in send buffer	[int MPI_Count	INTEGER
	datatype	datatype of elements of send buffer	MPI_Datatype	TYPE(MPI_Datatype)
	op	operation	MPI_Op	TYPE(MPI_Op)
	comm	communicator	MPI_Comm	TYPE(MPI_Comm)
	)			

```

1 template<typename T , typename F >
2 void mpl::communicator::allreduce
3   ( F, const T &, T & ) const;
4   ( F, const T *, T *,
5     const contiguous_layout< T > & ) const;
6   ( F, T & ) const;
7   ( F, T *, const contiguous_layout< T > & ) const;
8 F : reduction function
9 T : type

```

Allreduce arguments

Scalar buffer

```
1 double in_x,out_x;  
2 MPI_Allreduce( &in_x,&out_x,1,MPI_DOUBLE, ... )
```

Note: count= 1 because reduction on single scalar.

General buffer:

```
1 double in_x[5],out_x[5];  
2 MPI_Allreduce( in_x,out_x,5,MPI_DOUBLE, ... )
```

C++

```
1 vector<double> x(5);  
2 MPI_Comm_allreduce( x.data(), ...,5,... )
```

- *MPI_Datatype* is *MPI_INT*, *MPI_FLOAT* (for C), *MPI_REAL8* (for Fortran) et cetera.
- *MPI_Op* is *MPI_SUM*, *MPI_MAX* et cetera.

Allreduce arguments, MPL

Elementary types handled through overloading / templating.

```
1 float proc_data, reduce_data;  
2 comm_world.allreduce(mpl::plus<float>(), proc_data, reduce_data);  
3  
4 vector<float> x(2), y(2);  
5 mpl::contiguous_layout<float> two_floats(2);  
6 comm_world.allreduce  
7     (mpl::plus<float>(), x.data(), y.data(), two_floats);
```

Operators:

```
1 // separate recv buffer  
2 comm_world.allreduce(mpl::plus<float>(), proc_data, reduce_data);  
3 // in place  
4 comm_world.allreduce(mpl::plus<float>(), proc_data);
```

General type:

`T(T&)`

Available: `max`, `min`, `plus`, `multiplies`, `logical_and`, `logical_or`, `logical_xor`,
`bit_and`, `bit_or`, `bit_xor`.

Elementary datatypes

C	Fortran	Python	meaning
<i>MPI_CHAR</i>	<i>MPI_CHARACTER</i>	<i>MPI.DOUBLE</i>	only for text 8 bits like the C/F types
<i>MPI_SHORT</i>	<i>MPI_BYTE</i>		
<i>MPI_INT</i>	<i>MPI_INTEGER</i>		
<i>MPI_FLOAT</i>	<i>MPI_REAL</i>		
<i>MPI_DOUBLE</i>	<i>MPI_DOUBLE_PRECISION</i>		
	<i>MPI_COMPLEX</i>		
	<i>MPI_LOGICAL</i>		
unsigned	extensions		
			<i>MPI_Aint</i> <i>MPI_Offset</i>

A bunch more.

Reduction operators

MPI type	Python	meaning	applies to
MPI.Op			
<i>MPI_MAX</i>	<i>MPI.MAX</i>	maximum	integer, floating point
<i>MPI_MIN</i>	<i>MPI.MIN</i>	minimum	
<i>MPI_SUM</i>	<i>MPI.SUM</i>	sum	integer, floating point, complex, multilanguage types
<i>MPI_PROD</i>	<i>MPI.PROC</i>	product	
<i>MPI_REPLACE</i>	<i>MPI.REPLACE</i>	overwrite	C integer, logical
<i>MPI_NO_OP</i>	<i>MPI.OP_NULL</i>	no change	
<i>MPI_LAND</i>	<i>MPI.LAND</i>	logical and	
<i>MPI_LOR</i>	<i>MPI.LOR</i>	logical or	
<i>MPI_LXOR</i>	<i>MPI.LXOR</i>	logical xor	integer, byte, multilanguage types
<i>MPI_BAND</i>	<i>MPI.BAND</i>	bitwise and	
<i>MPI_BOR</i>	<i>MPI.BOR</i>	bitwise or	
<i>MPI_BXOR</i>	<i>MPI.BXOR</i>	bitwise xor	
<i>MPI_MAXLOC</i>	<i>MPI.MAXLOC</i>	max value and location	<i>MPI_DOUBLE_INT</i> and such
<i>MPI_MINLOC</i>	<i>MPI.MINLOC</i>	min value and location	

Exercise 4 randommax

Let each process compute a random number, and compute the sum of these numbers using the *MPI_Allreduce* routine.

$$\xi = \sum_i x_i$$

Each process then scales its value by this sum.

$$x'_i \leftarrow x_i / \xi$$

Compute the sum of the scaled numbers

$$\xi' = \sum_i x'_i$$

and check that it is 1.

Reduction to single process

Reduce with a single root process: great for printing out summary information at the end of your job.

Can you think of a case where a rooted reduce is appropriate or unavoidable? (Hint: tree)

Reduction to root

```
1 int MPI_Reduce
2   (void *sendbuf, void *recvbuf,
3    int count, MPI_Datatype datatype,
4    MPI_Op op, int root, MPI_Comm comm)
```

- Buffers: *sendbuf*, *recvbuf* are ordinary variables/arrays.
- Every process has data in its *sendbuf*,
Root combines it in *recvbuf* (ignored on non-root processes).
- *count* is number of items in the buffer: 1 for scalar.
- *MPI_Op* is *MPI_SUM*, *MPI_MAX* et cetera.

In-place operations

```
1 // allreduceinplace.c
2 for (int irand=0; irand<nrandoms; irand++)
3     myrandoms[irand] = (float) rand()/(float)RAND_MAX;
4 // add all the random variables together
5 MPI_Allreduce(MPI_IN_PLACE,myrandoms,
6               nrandoms,MPI_FLOAT,MPI_SUM,comm);
```

In-place (MPL)

Scalar:

```
1 // separate recv buffer
2 comm_world.allreduce(mpl::plus<float>(), proc_data, reduce_data);
3 // in place
4 comm_world.allreduce(mpl::plus<float>(), proc_data);
```

Buffer:

```
1 // collectbuffer.cxx
2 vector<float> rank2p2p1{ 2*xrank, 2*xrank+1 }, reduce2p2p1{0,0};
3 mpl::contiguous_layout<float> two_floats(rank2p2p1.size());
4 comm_world.allreduce
5   (mpl::plus<float>(), rank2p2p1.data(), reduce2p2p1.data(), two_floats);
6 if ( iprint )
7   cout << "Got: " << reduce2p2p1.at(0) << ", "
8       << reduce2p2p1.at(1) << endl;
```

Broadcast

```
1 int MPI_Bcast(  
2     void *buffer, int count, MPI_Datatype datatype,  
3     int root, MPI_Comm comm )
```

- All processes call with the same argument list
- *root* is the rank of the process doing the broadcast
- Each process allocates buffer space;
root explicitly fills in values,
all others receive values through broadcast call.
- Datatype is *MPI_FLOAT*, *MPI_INT* et cetera, different between C/Fortran.
- *comm* is usually *MPI_COMM_WORLD*

Gauss-Jordan elimination

<https://youtu.be/aQYuwat1WME>

MPI_Bcast

Name	Param name	Explanation	C type	F type
MPI_Bcast (				
MPI_Bcast_c (				
	buffer	starting address of buffer	void*	TYPE(*), DIMENSION(..)
	count	number of entries in buffer	[int MPI_Count	INTEGER
	datatype	datatype of buffer	MPI_Datatype	TYPE(MPI_Datatype)
	root	rank of broadcast root	int	INTEGER
	comm	communicator	MPI_Comm	TYPE(MPI_Comm)
	)			

```
1 template<typename T >
2 void mpl::communicator::bcast
3   ( int  root, T & data ) const
4   ( int  root, T * data, const layout< T > & l ) const
```

Exercise 5 jordan

The *Gauss-Jordan algorithm* for solving a linear system with a matrix A (or computing its inverse) runs as follows:

```
for pivot  $k = 1, \dots, n$ 
  let the vector of scalings  $\ell_i^{(k)} = A_{ik}/A_{kk}$ 
  for row  $r \neq k$ 
    for column  $c = 1, \dots, n$ 
       $A_{rc} \leftarrow A_{rc} - \ell_r^{(k)} A_{kc}$ 
```

where we ignore the update of the righthand side, or the formation of the inverse.

Let a matrix be distributed with each process storing one column. Implement the Gauss-Jordan algorithm as a series of broadcasts: in iteration k process k computes and broadcasts the scaling vector $\{\ell_i^{(k)}\}_i$. Replicate the right-hand side on all processors.

Exercise (optional) 6

Bonus exercise: can you extend your program to have multiple columns per process?

Scan

Scan

Scan or 'parallel prefix': reduction with partial results

- Useful for indexing operations:
- Each process has an array of n_p elements;
- My first element has global number $\sum_{q < p} n_q$.
- Two variants: *MPI_Scan* inclusive, and *MPI_Exscan* exclusive.

In vs Exclusive

process :	0	1	2	$\dots$	$p - 1$
data :	x_0	x_1	x_2	$\dots$	x_{p-1}
inclusive :	x_0	$x_0 \oplus x_1$	$x_0 \oplus x_1 \oplus x_2$	$\dots$	$\bigoplus_{i=0}^{p-1} x_i$
exclusive :	unchanged	x_0	$x_0 \oplus x_1$	$\dots$	$\bigoplus_{i=0}^{p-2} x_i$

MPI_Scan

Name	Param name	Explanation	C type	F type
MPI_Scan (				
MPI_Scan_c (				
	sendbuf	starting address of send buffer	const void*	TYPE(*), DIMENSION(..)
	recvbuf	starting address of receive buffer	void*	TYPE(*), DIMENSION(..)
	count	number of elements in input buffer	[int MPI_Count	INTEGER
	datatype	datatype of elements of input buffer	MPI_Datatype	TYPE(MPI_Datatype)
	op	operation	MPI_Op	TYPE(MPI_Op)
	comm	communicator	MPI_Comm	TYPE(MPI_Comm)
	)			

```

1 template<typename T , typename F >
2 void mpl::communicator::scan
3   ( F, const T &, T & ) const;
4   ( F, const T *, T *,
5     const contiguous_layout< T > & ) const;
6   ( F, T & ) const;
7   ( F, T *, const contiguous_layout< T > & ) const;
8 F : reduction function
9 T : type

```

MPI_Exscan

Name	Param name	Explanation	C type	F type
MPI_Exscan (				
MPI_Exscan_c (				
	sendbuf	starting address of send buffer	const void*	TYPE(*), DIMENSION(..)
	recvbuf	starting address of receive buffer	void*	TYPE(*), DIMENSION(..)
	count	number of elements in input buffer	[int MPI_Count	INTEGER
	datatype	datatype of elements of input buffer	MPI_Datatype	TYPE(MPI_Datatype)
	op	operation	MPI_Op	TYPE(MPI_Op)
	comm	intra-communicator	MPI_Comm	TYPE(MPI_Comm)
	)			

```

1 template<typename T , typename F >
2 void mpl::communicator::exscan
3   ( F, const T &, T & ) const;
4   ( F, const T *, T *,
5     const contiguous_layout< T > & ) const;
6   ( F, T & ) const;
7   ( F, T *, const contiguous_layout< T > & ) const;
8 F : reduction function
9 T : type


```

Exercise 7 scangather

- Let each process compute a random value n_{local} , and allocate an array of that length. Define

$$N = \sum n_{\text{local}}$$

- Fill the array with consecutive integers, so that all local arrays, laid end-to-end, contain the numbers $0 \cdots N - 1$. (See figure 7.)



Gather/Scatter, Barrier, and others

MPI_Gather

Name	Param name	Explanation	C type	F type
MPI_Gather (				
MPI_Gather_c (				
	sendbuf	starting address of send buffer	const void*	TYPE(*), DIMENSION(..)
	sendcount	number of elements in send buffer	[int MPI_Count	INTEGER
	sendtype	datatype of send buffer elements	MPI_Datatype	TYPE(MPI_Datatype)
	recvbuf	address of receive buffer	void*	TYPE(*), DIMENSION(..)
	recvcount	number of elements for any single receive	[int MPI_Count	INTEGER
	recvtype	datatype of recv buffer elements	MPI_Datatype	TYPE(MPI_Datatype)
	root	rank of receiving process	int	INTEGER
	comm	communicator	MPI_Comm	TYPE(MPI_Comm)
	)			

MPL: MPI_Gather

```
1 void mpl::communicator::gather
2   ( int  root_rank, const T * senddata, const layout< T > & sendl ) const
3   ( int  root_rank, const T * senddata, const layout< T > & sendl,
4     T * recvdata, const layout< T > & recvl ) const
5   // non-root versions:
6   ( int  root_rank, const T & senddata ) const
7   ( int  root_rank, const T & senddata, T * recvdata ) const
```

MPI_Scatter

Name	Param name	Explanation	C type	F type
MPI_Scatter (				
MPI_Scatter_c (				
	sendbuf	address of send buffer	const void*	TYPE(*), DIMENSION(..)
	sendcount	number of elements sent to each process	[int MPI_Count	INTEGER
	sendtype	datatype of send buffer elements	MPI_Datatype	TYPE(MPI_Datatype)
	recvbuf	address of receive buffer	void*	TYPE(*), DIMENSION(..)
	recvcount	number of elements in receive buffer	[int MPI_Count	INTEGER
	recvtype	datatype of receive buffer elements	MPI_Datatype	TYPE(MPI_Datatype)
	root	rank of sending process	int	INTEGER
	comm	communicator	MPI_Comm	TYPE(MPI_Comm)
)				


MPL: MPI_Scatter

```
1 void mpl::communicator::scatter
2   ( int  root_rank, const T *  send_data, const layout< T > &  sendl,
3     T *  recv_data, const layout< T > &  recvl ) const
4   ( int  root_rank, const T *  send_data,
5     T &  recv_data ) const
6   // non-root versions:
7   ( int  root_rank, T &  recv_data ) const
8   ( int  root_rank, T *  recv_data, const layout< T > &  recvl ) const
```

Gather/Scatter



- Compare buffers to reduce
- Scatter: the sendcount / Gather: the recvcount:
this is not, as you might expect, the total length of the buffer;
instead, it is the amount of data to/from each process.

Gather pictured



Popular application of gather

Matrix is constructed distributed, but needs to be brought to one process:



This is not efficient in time or space. Do this only when strictly necessary.
Remember SPMD: try to keep everything symmetrically parallel.

MPI_Allgather


Name	Param name	Explanation	C type	F type
MPI_Allgather (				
MPI_Allgather_c (				
	sendbuf	starting address of send buffer	const void*	TYPE(*), DIMENSION(..)
	sendcount	number of elements in send buffer	[int MPI_Count	INTEGER
	sendtype	datatype of send buffer elements	MPI_Datatype	TYPE(MPI_Datatype)
	recvbuf	address of receive buffer	void*	TYPE(*), DIMENSION(..)
	recvcount	number of elements received from any process	[int MPI_Count	INTEGER
	recvtype	datatype of receive buffer elements	MPI_Datatype	TYPE(MPI_Datatype)
	comm	communicator	MPI_Comm	TYPE(MPI_Comm)
	)			

```

1 void allgather
2   ( const T & send_data, T * recv_data ) const
3   ( const T * send_data, const layout< T > & sendl,
4     T * recv_data, const layout< T > & recvl ) const

```

Allgather pictured



V-type collectives

- Gather/scatter but with individual sizes
- Requires displacement in the gather/scatter buffer

MPI_Gatherv

Name	Param name	Explanation	C type	F type
MPI_Gatherv (				
MPI_Gatherv_c (				
	sendbuf	starting address of send buffer	const void*	TYPE(*), DIMENSION(..)
	sendcount	number of elements in send buffer	[int MPI_Count	INTEGER
	sendtype	datatype of send buffer elements	MPI_Datatype	TYPE(MPI_Datatype)
	recvbuf	address of receive buffer	void*	TYPE(*), DIMENSION(..)
	recvcunts	non-negative integer array (of length group size) containing the number of elements that are received from each process	[const int[] MPI_Count[]	INTEGER(*)
	displs	integer array (of length group size). Entry i specifies the displacement relative to recvbuf at which to place the incoming data from process i	[const int[] MPI_Aint[]	INTEGER(*)
	recvtype	datatype of recv buffer elements	MPI_Datatype	TYPE(MPI_Datatype)
	root	rank of receiving process	int	INTEGER
	comm	communicator	MPI_Comm	TYPE(MPI_Comm)
)				

MPL: MPI_Gatherv

```
1 template<typename T>
2 void gatherv
3     (int root_rank, const T *senddata, const layout<T> &sendl,
4      T *recvdata, const layouts<T> &recvls, const displacements &recvdispls) const
5     (int root_rank, const T *senddata, const layout<T> &sendl,
6      T *recvdata, const layouts<T> &recvls) const
7     (int root_rank, const T *senddata, const layout<T> &sendl ) const
```

Exercise 8 scangather

Take the code from exercise 7 and extend it to gather all local buffers onto rank zero. Since the local arrays are of differing lengths, this requires *MPI_Gatherv*.

How do you construct the lengths and displacements arrays?

Review 1


An *MPI_Scatter* call puts the same data on each process

```
/poll "A scatter call puts the same data on each process" "T" "F"
```

All-to-all

- Every process does a scatter;
- (equivalently: every process gather)
- each individual data, but amounts are identical
- Example: data transposition in FFT

Data transposition



Example: each process knows who to send to,
all-to-all gives information who to receive from

All-to-all

- Every process does a scatter or gather;
- each individual data and individual amounts.
- Example: radix sort by least-significant digit.

Radix sort

Sort 4 numbers on two processes:



	proc0		proc1	
array	2	5	7	1
binary	010	101	111	001
stage 1				
last digit	0	1	1	1
	(this serves as bin number)			
sorted	010		101	111 001
stage 2				
next digit	1		0	1 0
	(this serves as bin number)			
sorted	101	001	010	111
stage 3				
next digit	1	0	0	1
	(this serves as bin number)			
sorted	001	010	101	111
decimal	1	2	5	7

Reduce-scatter

- Pointwise reduction (one element per process) followed by scatter
- Somewhat related to all-to-all: data transpose but reduced information, rather than gathered.
- Applications in both sparse and dense matrix-vector product.

Example: sparse matrix setup

Example: each process knows who to send to,
all-to-all gives information how many messages to expect
reduce-scatter leaves only relevant information



Barrier

```
1 int MPI_Barrier( MPI_Comm comm )
```

- Synchronize processes:
- each process waits at the barrier until all processes have reached the barrier
- **This routine is almost never needed:**
collectives are already a barrier of sorts, two-sided communication is a local synchronization
- One conceivable use: timing

User-defined operators

MPI Operators

Define your own reduction operator

- Define operator between partial result and new operand

```
1 typedef void MPI_User_function
2   ( void *invec, void *inoutvec, int *len,
3     MPI_Datatype *datatype);
```

- Don't forget to free:

```
1 int MPI_Op_free(MPI_Op *op)
```

- Make your own reduction scheme *MPI_Reduce_local*

MPI_Op_create

Name	Param name	Explanation	C type	F type
MPI_Op_create (MPI_Op_create_c (user_fn commute op))				
	user_fn	user defined function	[MPI_User_function* MPI_User_function_c*	PROCEDURE (MPI_User_function)
	commute	true if commutative; false otherwise.	int	LOGICAL
	op	operation	MPI_Op*	TYPE(MPI_Op)

Missing MPL proto: mpi'op'create

Example

Smallest nonzero:

```
1 *(int*)inout = m;  
2 }
```

Review 2

The $\|\cdot\|_2$ norm (sum of squares) needs a custom operator.

/poll "The sum of squares norm needs a custom operators" "T" "F"

User-defined operators, MPL

General type:

$T(T\&)$

User-defined operator:

```
comm_world.reduce(lcm<int>(), 0, v, result);
```

Performance of collectives

Naive realization of collectives

Broadcast:



Single message:


$$\alpha = \text{message startup} \approx 10^{-6}s, \quad \beta = \text{time per word} \approx 10^{-9}s$$

- Time for message of n words:

$$\alpha + \beta n$$

- Time for collective? Can you improve on that?

Better implementation of collective



- What is the running time now?
- Can you come up with lower bounds on the α, β terms? Are these achieved here?
- How about the case of really long buffers?

Implementation of Reduce

	$t = 1$	$t = 2$	$t = 3$
p_0	$x_0^{(0)}, x_1^{(0)}, x_2^{(0)}, x_3^{(0)}$	$x_0^{(0:1)}, x_1^{(0:1)}, x_2^{(0:1)}, x_3^{(0:1)}$	$x_0^{(0:3)}, x_1^{(0:3)}, x_2^{(0:3)}, x_3^{(0:3)}$
p_1	$x_0^{(1)} \uparrow, x_1^{(1)} \uparrow, x_2^{(1)} \uparrow, x_3^{(1)} \uparrow$		
p_2	$x_0^{(2)}, x_1^{(2)}, x_2^{(2)}, x_3^{(2)}$	$x_0^{(2:3)} \uparrow, x_1^{(2:3)} \uparrow, x_2^{(2:3)} \uparrow, x_3^{(2:3)} \uparrow$	
p_3	$x_0^{(3)} \uparrow, x_1^{(3)} \uparrow, x_2^{(3)} \uparrow, x_3^{(3)} \uparrow$		

Implementation of Allreduce

	$t = 1$	$t = 2$	$t = 3$
p_0	$x_0^{(0)} \downarrow, x_1^{(0)} \downarrow, x_2^{(0)} \downarrow, x_3^{(0)} \downarrow$	$x_0^{(0:1)} \downarrow\downarrow, x_1^{(0:1)} \downarrow\downarrow, x_2^{(0:1)} \downarrow\downarrow, x_3^{(0:1)} \downarrow\downarrow$	$x_0^{(0:3)}, x_1^{(0:3)}, x_2^{(0:3)}, x_3^{(0:3)}$
p_1	$x_0^{(1)} \uparrow, x_1^{(1)} \uparrow, x_2^{(1)} \uparrow, x_3^{(1)} \uparrow$	$x_0^{(0:1)} \downarrow\downarrow, x_1^{(0:1)} \downarrow\downarrow, x_2^{(0:1)} \downarrow\downarrow, x_3^{(0:1)} \downarrow\downarrow$	$x_0^{(0:3)}, x_1^{(0:3)}, x_2^{(0:3)}, x_3^{(0:3)}$
p_2	$x_0^{(2)} \downarrow, x_1^{(2)} \downarrow, x_2^{(2)} \downarrow, x_3^{(2)} \downarrow$	$x_0^{(2:3)} \uparrow\uparrow, x_1^{(2:3)} \uparrow\uparrow, x_2^{(2:3)} \uparrow\uparrow, x_3^{(2:3)} \uparrow\uparrow$	$x_0^{(0:3)}, x_1^{(0:3)}, x_2^{(0:3)}, x_3^{(0:3)}$
p_3	$x_0^{(3)} \uparrow, x_1^{(3)} \uparrow, x_2^{(3)} \uparrow, x_3^{(3)} \uparrow$	$x_0^{(2:3)} \uparrow\uparrow, x_1^{(2:3)} \uparrow\uparrow, x_2^{(2:3)} \uparrow\uparrow, x_3^{(2:3)} \uparrow\uparrow$	$x_0^{(0:3)}, x_1^{(0:3)}, x_2^{(0:3)}, x_3^{(0:3)}$

Review 3

True or false: there are collectives that do not communicate data

/poll "there are collectives that do not communicate data" "T" "F"

Reduction operators

User-defined operators

Given a reduction function:

```
1 typedef void user_function
2   ( void *invec, void *inoutvec, int *len,
3     MPI_Datatype *datatype);
```

create a new operator:

```
1 MPI_Op rwz;
2 MPI_Op_create(user_function,1,&rwz);
3 MPI_Allreduce(data+procno,&positive_minimum,1,MPI_INT,rwz,comm);
```

Exercise 9 onenorm

Write the reduction function to implement the *one-norm* of a vector:

$$\|x\|_1 \equiv \sum_i |x_i|.$$