

CUDA Course

Victor Eijkhout

2025 COE 379L / CSE 392

Justification


CUDA is the basic programming model for NVidia GPUs ...
... which TACC is getting a lot of.

1. Introduction

Why GPUs

- Massive parallelism (but that's really multicore and SIMD)
- Simpler processor: less power / more computing for the same power.
- Simpler processor: no general programs, only 'data parallel' operations
... such as machine learning
- Latest GPUs have very high speed on low precision
... such as for machine learning

Attached processor



- GPU is attached to regular CPU
- computation is 'offloaded'
- GPU has its own memory:
you may need to synchronize

Data parallel programming

- Kernel: operation for a single thread;
in effect: body of inner loop
- Grid: multi-dimensional structure of available 'threads'
executed on cores
(not a CPU core, more SIMD vector lane)

Processor structure

- Streaming Multiprocessor: like a core
- Warp: vector instruction
- Thread: like a simd-lane
- Single Instruction Multiple Thread (SIMT) 'massive multi-threading':
you can have way more 'threads' than $\text{blocks} \times \text{blocksize}$
processor can switch very fast between threads

Loops vs CUDA

```
1 float *x;
2 for ( blockIdx_x < gridsize )
3     for ( threadIdx_x < blocksize )
4         size_t linear = threadIdx_x
5             + blockIdx_x * blocksize;
6         f( x[linear] );
```

```
1 __global__ cu_f( float x) {
2     size_t linear = threadIdx_x
3         + blockIdx_x * blocksize;
4     f( x[linear] );
5 }
6
7 float *x; // on device
8 cu_f<<<gridsize,blocksize>>>(x);
```


GPU memory

- Very much graphics inspired:
 'constant' memory, 'texture' memory
- Main memory: has grown from 8G to 100G-ish
- Shared memory inside each SM: managed cache
- Actual caches and registers.

Comparison

- Dozens of cores
- Core has vector instructions
- a vector instruction has 'lanes'
- Latency hiding by caches, prefetching
- Dozens of 'Streaming Multiprocessors'
- SM has 'warp's
- a warp has threads
- Latency hiding by massive multi-threading


GPU cluster

Memory GPU

GPU

CPU

Memory CPU



Step 1:
One node


Step 2:
Multiple nodes

- Use GPU on a node
(possible combined with OpenMP or other threading)
- Use MPI between nodes;
NVlink also exists but not on large scale

TACC's Vista


- 256 Grace-Grace two-socket CPU nodes
4 racks of 64 nodes each
- Grace-grace node has twice 72 cores: 144 total per node.
- 600 Grace-Hopper nodes
19 racks of 32 nodes
- GH200 Grace-Hopper superchip combines
each containing one 72-core Grace CPU (120 GiB LPDDR5X)
and H200 GPU (96 GiB HBM3)
- GH node: 34 Tflops of FP64
1979 Tflops of FP16 for Machine Learning.

Grace Hopper superchip



Grace Hopper superchip

Grace-Hopper GPU/CPU Memory Coherence



single page table enabled by address translation services (ATS)



71

- Single page table
- CPU and GPU can access data by paging.



CUDA and others

- CUDA is for C and Fortran;
limited to NVidia
- OpenACC, OpenCL: portable but tough to program
- AMD GPUs: 'HIP'
- Intel GPUs (hm): Sycl
- OpenMP offload: universal but maybe not highest performance
- Kokkos: portable to OpenMP and all GPUs

2. Kernels

Starting a kernel

Ordinary C/C++ function;

`__global__` means executable on device:

```
1 // hello.cu
2 __global__ void helloWorldKernel() {
3     printf("Hello, World!\n");
4 }
```

Execute on device with blocksize and gridsize
(here both one, so one core):

```
1 helloWorldKernel<<<1, 1>>>();
```

Translation code to kernel

```
1 for ( int i=0; i<n; ++i ) {
2     int threadIdx.x = i;
3     f(i);
4 }
5
6 __global__ computef( /* fargs */ ) {
7     int i = threadIdx.x;
8     f(i);
9 }
10
11 int gridsize,blocksize;
12 computef<<<gridsize,blocksize>>>>( fargs );
```

- Kernel is the body of the inner loop;
- Kernel is executed for each thread in the grid;
- Kernel consults `threadIdx` to find the 'iteration' number.
- Note: kernel invocation returns immediately, use `cudaDeviceSynchronize` to wait until device finishes

About threads

- CUDA creates thread blocks until all data has been processed
- You can tune the block size up to some limit;
- Make sure to create enough blocks.

Code:

```
1 // blocksizes.cu
2 printf("Device Name: %s\n", prop.name);
3 printf("Max threads per block: %d\n",
4        prop.maxThreadsPerBlock);
5 printf("Max block dimensions:\n %d x %d x
6        %d\n",
7        prop.maxThreadsDim[0],
8        prop.maxThreadsDim[1],
9        prop.maxThreadsDim[2]);
```

Output:


```
1 Device Name: NVIDIA
   ↪ GH200 120GB
2 Max threads per block:
   ↪ 1024
3 Max block dimensions:
4 1024 x 1024 x 64
```

Thread indexing

```
1 int global_thread_id = blockIdx.x * blockDim.x + threadIdx.x; // or 2d, 3d variant
2 if ( global_thread_id >= datasize ) return;
```

- Global thread index from block and thread coordinate;
- Global index often used as index in the data;
- You can have more threads than data;
no problem but you do need to test.


Thread indexing illustrated



Thread indexing rationale

Why consecutive threads on consecutive array locations?

Why not let each thread have a consecutive chunk?



Data on the device

```
1 float
2  *x, *y, *z;
3 size_t nbytes = N*sizeof(float);
4 CU_ASSERT( cudaMallocManaged( &x,nbytes ) );
5 CU_ASSERT( cudaMallocManaged( &y,nbytes ) );
6 CU_ASSERT( cudaMallocManaged( &z,nbytes ) );
7
8 // Run kernel on 1M elements on the CPU
9 int blocksize = 1024;
10 int gridsize = N/blocksize+1;
11 add1<<<gridsize,blocksize>>>(N,x,y,z);
12 // Wait for GPU to finish before accessing data on host
13 CU_ASSERT( cudaDeviceSynchronize() );
```

- Managed allocation is accessible on host and device
- Device is asynchronous, so synchronization needed

Utility macro

```
1 #define CU_ASSERT( code ) { \  
2     cudaError_t err = code ; \  
3     if (err!=cudaSuccess) { \  
4         printf("error <<%s>> on line %d\n",cudaGetErrorString(err),__LINE__); \  
5         return 1; } }
```

- CUDA routines can return error code
- Check on fatal errors.

Exercise 1

Write a CUDA kernel that adds two arrays into a third.

Using the above snippets, finish the code, test for correctness.

Two-dimensional grid

```
1 dim3 blocksize(16,16);
2 dim3 gridsize( rowsize/blocksize.x+1, colsize/blocksize.y+1 );
3 printf(" .. grid is (%d,%d) blocks of (%d,%d)\n",
4         gridsize.x, gridsize.y, blocksize.x, blocksize.y);
```

- Two-dimensional data calls for two-dimensional thread organization (same for three of course)
- Specify grid/block size with `dim3` objects.

Two-dimensional indexing

```
1 // index2d.cu
2 // Calculate the row and column indices for this thread
3 size_t row = blockIdx.y * blockDim.y + threadIdx.y;
4 size_t col = blockIdx.x * blockDim.x + threadIdx.x;
5
6 // Make sure we don't go out of bounds
7 if (row < height && col < width) {
8 // Calculate the linear index for the matrices
9 size_t idx = row * width + col;
10
11 // Perform the addition
12 d_c[idx] = d_a[idx] + d_b[idx];
```

- Compute row/col from x/y components
- Adjacent threads need to access adjacent data

Thread numbering

0	1	2
3	4	5

6	7	
		11

12		
		17

18		
		23

- Lexicographic ordering of blocks;
- Lexicographic ordering of threads in blocks.

Exercise 2

Create a two-dimensional array and let each thread write its global number in the element where it is active:


```
1 output[output_loc] = global_thread_num;
```

Your program should recreate the layout of figure 28.

3. Memory

Host and device memory

Grace/Hopper Memory Allocation Types



TACC

72

- Host memory can be copied;
- can be visible on the device;
- can be unified with the device.

Explicit transfer

Allocate input data and copy host values:

```
1 // input allocated on device
2 int *dev_input;
3 cudaMalloc( &dev_input, bytesize );
4 // host values copied
5 cudaMemcpy( dev_input, input, bytesize, cudaMemcpyHostToDevice );
```

After the kernel execution copy output data back:

```
1 cudaDeviceSynchronize();
2 cudaMemcpy( output, dev_output, bytesize, cudaMemcpyDeviceToHost );
```

Free:

```
1 free(input); free(output);
2 cudaFree(dev_input); cudaFree(dev_output);
```


- Host holds pointer to device memory;
- used for copying in/out;
- device memory freed.

Implicit transfer

```
1 // add1m.cu
2 size_t nbytes = N*sizeof(float);
3 CU_ASSERT( cudaMallocManaged( &x,nbytes ) );
4 CU_ASSERT( cudaMallocManaged( &y,nbytes ) );
5 CU_ASSERT( cudaMallocManaged( &z,nbytes ) );
```

- Possibility one:
allocate on the host with `cudaHostAlloc`
using a flag of `cudaHostAllocMapped`.
Pass `cudaHostGetDevicePointer` to the device.
- Easier: `cudaMallocManaged`;
transfer is done through on-demand paging.
- (Implications for timing!)

GPU memory



Shared memory

```
1 __shared__ float [32] [32];
```

- Shared between threads in a block:
managed cache per core
- Can be used as fast local storage
- Copy from RAM to 'shared', operate on shared, copy back.

4. Synchronization

Mechanisms

- Threads in a warp are always synchronized.
- Threads in a block can be synchronized with `__syncthreads`.
- Blocks in a grid can not be synchronized.
- Host and device can be synchronized with `cudaDeviceSynchronize`.

Warps and warp divergence

```
1 if ( threadIdx.%2==0 )  
2   functionA();  
3 else  
4   functionB();
```

```
1 bool mask = (threadIdx.x%2==0)  
2 if (mask) functionA();  
3 if (not mask) functionB();
```

- Threads in a warp are synchronized
- Conditionals are serialized: *warp divergence*
- Try to rewrite your code...

Synchronization across blocks

- Threads blocks are not synchronized; can not be.
Does this make sense from the original purpose of Graphics Processing Units (GPUs)?
- Solution 1: synchronize outside CUDA;
- Solution 2: use atomics;
- Solution 3: use a library that does it for you.

Example: sum reduction


- Recursive summation in each block:

```
1 // redumy.cu
2 __global__
3 void reduce0(size_t n, float *input, float *y);
```

- Summing the blocks outside of CUDA:

```
1 for ( int b=0; b<gridsize; ++b)
2   value += y[b*blocksize];
```


Implementation 1




We let threads at distances 2, 4, 8, ... sum values from offsets 1, 2, 4, ...

Code

```
1 __global__
2 void reduce0(size_t n, float *input, float *y) {
3     size_t
4         thrd = threadIdx.x,
5         global = blockDim.x*blockIdx.x + threadIdx.x;
6     if (global>=n) return;
7     y[global] = input[global];
8     __syncthreads();
9     for ( size_t offset=1; offset<=blockDim.x/2; offset*=2 ) {
10         if (thrd%(2*offset)==0)
11             y[global] += y[global+offset];
12         __syncthreads();
13     }
14 }
```

Implementation 2



- Non-obvious mapping thread $\leftrightarrow$ data

Code

```
1 // pointer to this block
2 float *data = y+blockDim.x*blockIdx.x;
3 for ( size_t offset=1; offset<=blockDim.x/2; offset*=2 ) {
4     size_t index = 2 * offset * thrd;
5     if (index+offset<blockDim.x)
6         data[index] += data[index+offset];
7     __syncthreads();
8 }
9 // we need the following statement really only for thread 0:
10 y[global] = data[thrd];
```

Implementation 3


- Use 'Cuda UnBound'
- <https://github.com/NVIDIA/cccl/tree/main/cub>

```
1 // compute needed tmp space
2 cub::DeviceReduce::Reduce
3   ( d_temp_storage, temp_storage_bytes,
4     x, &d_nrm, num_items, devplus{}, 0.f );
5 // alloc tmp space
6 cudaMalloc(&d_temp_storage, temp_storage_bytes);
7 // actual reduction
8 cub::DeviceReduce::Reduce
9   ( d_temp_storage, temp_storage_bytes,
10    x, &d_nrm, num_items, devplus{}, 0.f );
11 CU_ASSERT( cudaDeviceSynchronize() );
12 cudaFree(&d_temp_storage);
```

5. Example: stencil

Five-point stencil

$$y_{ij} \leftarrow 4x_{ij} - x_{i-1,j} - x_{i+1,j} - x_{i,j-1} - x_{i,j+1}$$



- Streaming kernel, but rhs elements used multiple times
- We use bordered data for easy of programming.

Coordinates

```
1 // smooth2d.cu
2 auto global_row = [=] ( int64_t by,int64_t ty ) {
3     return by * blockDim.y + ty; };
4 auto global_col = [=] ( int64_t bx,int64_t tx ) {
5     return bx * blockDim.x + tx; };
6 auto global_loc = [nsize] ( int64_t row,int64_t col ) {
7 // one blank row, and each row is left & right bordered
8     return (row+1) * (nsize+2) + col+1; };
9
10 int64_t bx = blockIdx.x, by = blockIdx.y, tx = threadIdx.x, ty = threadIdx.y;
11 int64_t my_global_row = global_row(by,ty);
12 int64_t my_global_col = global_col(bx,tx);
13 if ( my_global_row<0 or my_global_col<0
14     or my_global_row>=msize or my_global_col>=nsize )
15     return;
16
17 y[my_global_loc] = 4*x[my_global_loc]
18 - x[ global_loc( my_global_row-1,my_global_col ) ]
19 - x[ global_loc( my_global_row+1,my_global_col ) ]
20 - x[ global_loc( my_global_row ,my_global_col-1 ) ]
21 - x[ global_loc( my_global_row ,my_global_col+1 ) ];
```

- Utility functions,
- translation thread to global coordinate

Use of local memory

If elements read multiple times:


- Copy to local ('shared') data,
- Operate on that
- Write back

Shared memory

```
1 __shared__ float sums[BLOCKSIZE];  
2 sums[thrd_idx] = x[global_loc];  
3 __syncthreads();
```

- Shared memory size of thread block:
static dimensions
- Shared between all threads of block

Shared memory in stencil computation



- Assume bordered global domain,
- Make local block bordered too
- Tricky programming!

Exercise 3

Take your 2D vector addition code;
use shared memory. Do you see a performance difference?
If not, can you find a way to make this visible?

For instance, try:

```
1 __global__ void add_with_coef( double *out,double *in,  
2     double *coef,int ncoeff) {  
3     for ( icoeff=0; icoeff<ncoeff; ++icoeff ) {  
4         for ( .... i ... )  
5             a[i] += coef[icoeff] * in[i];  
6     }  
7 }
```

6. Timing, profiling and such

Device queries

```
1 // device.cu
2 int ndev;
3 auto status = cudaGetDeviceCount(&ndev);
4
5 cudaDeviceProp properties;
6 cudaGetDeviceProperties(&properties, iddev);
7
8 cout << "Device " << iddev << " = " << properties.name << '\n';
9 cout << "  async: " << properties.asyncEngineCount << '\n';
10 cout << "  unified: " << properties.unifiedAddressing << '\n';
11
12 cout << "  capability: " << properties.major << "." << properties.minor << '\n';
13 cout << "  multiprocs: " << properties.multiProcessorCount << '\n';
14 cout << "  clock rate: " << properties.clockRate << '\n';
15
16 cout << "  global memory: " << properties.totalGlobalMem << '\n';
17 cout << "  shared mem/block: " << properties.sharedMemPerBlock << '\n';
```

Block size

Code:

```
1 // blocksize.cu
2 printf("Device Name: %s\n", prop.name);
3 printf("Max threads per block: %d\n",
4        prop.maxThreadsPerBlock);
5 printf("Max block dimensions:\n %d x %d x
6        %d\n",
7        prop.maxThreadsDim[0],
8        prop.maxThreadsDim[1],
9        prop.maxThreadsDim[2]);
```

Output:

```
1 Device Name: NVIDIA
   ↪ GH200 120GB
2 Max threads per block:
   ↪ 1024
3 Max block dimensions:
4 1024 x 1024 x 64
```

```
1 cudaEvent_t start, stop;
2 cudaEventCreate(&start);
3 cudaEventCreate(&stop);
4 float milliseconds = 0.0f;
5
6 cudaEventRecord(start);
7 cudaEventSynchronize(start);
8 // device operation goes here
9 cudaEventRecord(stop);
10 cudaEventSynchronize(stop);
11 cudaEventElapsedTime(&milliseconds, start, stop);
```

- You can time on the host, but make sure to synchronize
- Time on the device with events, act like barriers.





Exercise 4 code